

images/uel.png

SCHOOL OF ARCHITECTURE, COMPUTING AND ENGINEERING
Department of Computing and Digital Technologies

Deepfake Detection Using Deep Learning Models

Jasmi Preethi Alasapuri
2571395

A dissertation submitted in partial fulfilment of the degree of
BSc (Hons) in Data Science and Artificial Intelligence.

Supervisor: Lucian Duta

CN6000

April 14, 2026

Abstract

This is the abstract. Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

Acknowledgements

These are the acknowledgements. Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

Contents

Abstract	ii
Contents	iv
List of Figures	vi
List of Tables	vii
1 Introduction	1
1.1 Sub-heading	1
2 Literature Review	2
2.1 Introduction	2
2.1.1 Define Deep Learning	2
2.1.2 Define Deepfake	3
2.2 Evolution of Deepfakes	3
2.2.1 Phase 1: Visual Fidelity and Generative Advances	3
2.2.2 Phase 2: “In-the-Wild” Adaptation and Multi-Subject Expansion	4
2.2.3 Phase 3: Multimodal and Neural Generation	4
2.3 Documented Real-World Impacts and High Profile Cases	5
2.4 The Dual Dilemma: Creative Application Versus Malicious Threat	7
2.5 Deepfake Generation Techniques	8
2.5.1 Image-Based Manipulation	8
2.5.2 Audio-Based Manipulation	8
2.5.3 Video-Based Manipulation	9
2.6 Deepfake Detection Methodologies	10
2.6.1 Image-Based Detection Approaches	10
2.6.2 Audio-Based Detection Approaches	10
2.6.3 Video Based Detection Approaches	11
2.6.4 Multimodal Detection Approaches	11
2.6.5 Generalisation and Cross Dataset Performance	12
2.7 Gaps in the Literature	12
2.8 Conclusion	13

3	Methodology	14
3.1	Introduction	14
3.2	Research and Development Approach	14
3.3	System Architecture and implementation	15
3.3.1	Overview of the Proposed System	15
3.3.2	Dataset Selection and Subset Strategy	15
3.3.3	Audio Feature Extraction Module	15
3.3.4	Visual Feature Extraction Module	15
3.3.5	Cross-Modal Fusion and Classification	16
3.4	Model Training and Evaluation Strategy	16
3.5	Challenges and Limitations	16
3.6	Ethical Considerations	17
3.7	Conclusion	17
4	Results/Findings	18
4.1	Source Code	18
5	Evaluation	19
6	Conclusions	20
	Appendix A Media Content	21
	Appendix B Installation Instructions	22
	Appendix C User Manual	23
	References	25

List of Figures

List of Tables

List of Abbreviations

Chapter 1: Introduction

Write you introduction here.

1.1 Sub-heading

This is an example of a subheading. [delete as appropriate] Typical headings include 1.1. Introduction, 1.2. Purpose of this project, 1.3. Problem Statement, 1.4. Aim of the project, 1.5. Objectives of the project, 1.6. Structure of the project.

Chapter 2: Literature Review

2.1 Introduction

Recent advances in artificial intelligence have reshaped digital media forensics, largely due to the rapid rise of deepfakes. These synthetic videos and images, generated through deep neural architectures, can imitate human appearance and behaviour with increasing precision (He et al., 2021; Rossler et al., 2019). Although early work treated manipulated media as a relatively contained technical issue, current developments have made detection more difficult. As He et al. (2021) pointed out, the realism achieved by modern generative models disrupts long-standing assumptions about what constitutes trustworthy audio-visual evidence. Rossler et al. (2019) further highlight that detection systems often are behind the pace of generative model improvements, creating a persistent asymmetry between manipulation and forensics. Given this trajectory, a clearer understanding of the underlying deep learning methods is necessary before evaluating why current detection strategies struggle and where they fall short.

2.1.1 Define Deep Learning

Deep learning is defined as a subset of machine learning characterized by the utilisation of multi-layered neural networks. This architectural paradigm draws inspiration from the biological neural networks of the human brain, aiming to model high-level abstractions in data through a deep graph with multiple processing layers (Lecun et al., 2015). According to Goodfellow et al. (2016), deep learning systems differentiate themselves from conventional machine learning by their ability to learn representations of data with multiple levels of abstraction. While a shallow network might only have one hidden layer, a "deep" architecture utilises a hierarchy of layers to process data through a series of nested nonlinear equations. This structure allows the model to function as a universal approximator, capable of modeling complex high-dimensional data patterns that simpler algorithms cannot capture (Hornik et al., 1989).

2.1.2 Define Deepfake

The term deepfake combines “deep learning” and “fake,” which describes synthetic media generated through deep neural models rather than conventional editing techniques (Westerlund, 2019). Most early systems relied on autoencoders or Generative Adversarial Networks, which made it possible to automate face swapping and expression transfer in a way that traditional manipulation tools could not achieve (Güera and Delp, 2018). Although digital alterations long predate deep learning, the specific phenomenon of deepfakes entered mainstream attention in late 2017 when a Reddit user released code that enabled realistic face swaps with very little technical skill, initially leading to widespread non consensual pornography and, later, to wider use in misinformation and entertainment (Chesney and Citron, 2019; Westerlund, 2019).

In research contexts, deepfakes are generally defined by their use of deep generative models to produce content that is difficult to distinguish from authentic footage Cai et al. (2024). Over time, beyond face swaps, current systems include audio-based voice cloning and multimodal models capable of synchronising generated speech with fabricated lip movements, which introduces additional complexity for forensic detection (Yi et al., 2023). Together, these developments show that deepfakes are not a single technique but a growing family of generative methods with expanding forensic implications.

2.2 Evolution of Deepfakes

2.2.1 Phase 1: Visual Fidelity and Generative Advances

The initial phase of deepfake development focused on face swapping, using simple autoencoders and early Generative Adversarial Networks (GANs) (Dolhansky et al., 2020). These techniques were prone to noticeable flaws, such as affine warping errors, making them easily detectable by early forensic tools (Li and Lyu, 2018). However, the field progressed rapidly with the adoption of improved training strategies such as the Two-Time Scale Update Rule (TTUR), which greatly improved GAN stability and convergence, resulting in more lifelike image generation (Heusel et al., 2017).

As generative methods evolved, benchmark datasets such as the Deepfake Detection Challenge (DFDC) (Dolhansky et al., 2020) and Celeb-DF (Li et al., 2019) emerged. These resources marked a transition from low-grade and easily spotted fakes to high-

quality videos that accurately replicated real-world visuals, posing significant challenges for detection systems (Li et al., 2019).

2.2.2 Phase 2: “In-the-Wild” Adaptation and Multi-Subject Expansion

A major advancement in deepfake technology was transitioning from controlled laboratory settings to more unpredictable “in-the-wild” environments. Although earlier datasets typically included individual subjects under uniform lighting and framing, newer deepfake datasets like Wild Deepfake (Zi et al., 2021) and FFIW-10K (Zhou et al., 2021) consist of content sourced from the internet or specifically generated to endure real-world challenges such as compression and noise.

During this stage, deepfake methods were also advanced, enabling the manipulation of multiple people within the same scene. For example, research by Narayan et al. (2023) using the DF-Platter data set showed that current techniques can alter multiple faces in a single image, even in challenging scenarios such as occlusion or poor image quality. Similarly, the FFIW-10K dataset highlighted the ongoing difficulty that forensic systems face in detecting fakes in group videos—especially when only some faces have been manipulated (Zhou et al., 2021).

2.2.3 Phase 3: Multimodal and Neural Generation

The latest and arguably most concerning development in deepfake technology is the move from only visual edits to multimodal deepfakes that manipulate both audio and video. Recent studies highlight the importance of inconsistencies between sound and visuals, leading to the creation of data sets such as AV-Deepfake1M (Cai et al., 2024) and FakeAVCeleb (Yi et al., 2023), which feature coordinated alterations to both video and audio streams. This evolution requires that detection tools now evaluate the synchronization of lip movements with speech, rather than just looking for visual inconsistencies (Yi et al., 2023).

In addition, the generative techniques themselves have progressed beyond traditional GANs. Neural Radiance Fields (NeRF) have been used to create talking-head videos where both viewpoint and audio can be modified Guo et al. (2021). For audio, generation has advanced from simple concatenation methods to diffusion models, such

as NaturalSpeech 2, which enable zero-shot voice synthesis with remarkably accurate intonation and prosody (Shen et al., 2023).

2.3 Documented Real-World Impacts and High Profile Cases

Although early research on deepfakes largely focused on improving synthesis quality and benchmarking detection accuracy, larger concern has been driven by their increasing use in real-world scenarios (Chesney and Citron, 2019; Dolhansky et al., 2020; Rossler et al., 2019).

One area of significant impact has been medical misinformation. Investigative reporting has documented numerous cases in which deepfake videos impersonating qualified doctors were used to promote false health advice and fraudulent treatments across social media platforms (Clark, 2025; ITVNews, 2024). These videos often exploit the visual credibility of trusted medical professionals, making misinformation difficult for non-expert audiences to identify. Such cases highlight how deepfakes can undermine public trust in expert knowledge, particularly in domains where inaccurate information may cause physical harm (Chesney and Citron, 2019; Rossler et al., 2019).

Deepfakes have also been employed in financial fraud and impersonation scams. A notable case involved a multinational engineering firm in which employees were deceived into authorising a large financial transfer after participating in a video call that appeared to feature senior executives but was later confirmed to be AI generated (Milmo, 2024). Similar incidents have been reported globally, prompting warnings from law-enforcement agencies about the growing use of synthetic audio and video in social engineering attacks (Bragg, 2025). These cases demonstrate how deepfakes can amplify traditional fraud techniques by adding a convincing visual and auditory layer.

In addition to visual manipulation, audio deepfakes have demonstrated significant real world impact, particularly in the context of fraud and impersonation (Korshunov and Marcel, 2018; Yi et al., 2023; Zhang et al., 2025). One of the first high profile cases occurred in 2019, when attackers used a synthetic voice mimicked a company executive to authorise a fraudulent financial transfer, resulting in substantial financial losses (Stupp, 2019). Unlike visual deepfakes, audio-based impersonation can be deployed in real time through phone calls or voice messages, limiting opportunities for human verifi-

cation and increasing its effectiveness in social engineering scenarios (Korshunov and Marcel, 2018). Prior research demonstrates that synthetic speech can convincingly replicate speaker identity, enabling attacks that bypass traditional speaker verification and authentication systems (Korshunov and Marcel, 2018; Yi et al., 2023). Taken together, these findings indicate that audio deepfakes pose a parallel threat to trust and identity assurance mechanisms, reinforcing the need for detection approaches that extend beyond visual analysis and address multimodal manipulation (Cai et al., 2024; Rossler et al., 2019).

Political manipulation represents another area of concern associated with deepfakes, particularly in the context of increased geopolitical tension (Chesney and Citron, 2019; Vaccari and Chadwick, 2020). Videos depicting public figures have circulated online, sometimes appearing to show officials making false statements or announcements, thereby creating confusion regarding the authenticity of political communication (Rossler et al., 2019). An example involves a manipulated video of a national leader that was briefly disseminated on social media before being removed by platform moderators (Allyn, 2022). Although such videos do not always achieve their intended persuasive impact, studies suggest that their rapid circulation can take advantage of periods of uncertainty and contribute to decline of confidence in authentic news sources (Vaccari and Chadwick, 2020). Deepfakes targeting individuals, ethical issues In addition to large-scale misinformation, deepfakes have caused direct personal harm. Investigations of non-consensual synthetic pornography have revealed extensive platforms dedicated to generating explicit content using the likenesses of individuals without consent, disproportionately targeting women (Moore, 2025). These cases underscore the ethical and legal challenges posed by deepfakes, particularly in situations where existing regulatory and legal frameworks struggle to address harms arising from fabricated yet highly realistic media (Chesney and Citron, 2019; Westerlund, 2019). Taken together, these examples illustrate that deepfakes are no longer a purely technical problem confined to academic datasets. Their real world deployment exposes weaknesses in current detection systems and reinforces the need for reliable, generalisable forensic approaches capable of operating under the unpredictable conditions of online media ecosystems (Dolhansky et al., 2020; Rossler et al., 2019).

2.4 The Dual Dilemma: Creative Application Versus Malicious Threat

Deepfake technologies exemplify a dual-use dilemma, in which the same deep generative models that enable legitimate and creative applications can also be exploited for harmful purposes. Advances in audio and visual synthesis have supported a range of beneficial use cases, including film post production, automated dubbing, digital de-ageing, and assistive technologies such as personalised text-to-speech systems and voice restoration tools. In these contexts, synthetic media techniques are typically deployed with consent and transparency, and their use is framed as an extension of existing digital production methods (Westerlund, 2019; Yi et al., 2023).

However, the properties that make deepfake technologies attractive for creative and commercial applications, namely high realism, automation, and scalability, also enable malicious misuse. As discussed in Section 2.3, deepfakes have been employed in misinformation campaigns, fraud, impersonation, and non consensual content generation. The ability to finetune generative models using limited data further reduces the barrier to entry for misuse, allowing individuals without advanced technical expertise to produce convincing synthetic media (Chesney and Citron, 2019).

This tension complicates regulatory and technical responses to deepfakes, as restrictive controls on generative technologies risk constraining legitimate innovation, while inadequate oversight allows harmful applications to proliferate (Chesney and Citron, 2019; Westerlund, 2019). From a forensic perspective, this dual use nature highlights the limitations of prevention only strategies and reinforces the need for effective detection mechanisms (Chesney and Citron, 2019). Detection systems must therefore be capable of operating independently of the intent behind content creation, instead focusing on identifying artefacts, inconsistencies, or statistical traces that distinguish synthetic media from authentic recordings (Dolhansky et al., 2020; Rossler et al., 2019).

Understanding deepfakes as a dual use technology provides an important context for evaluating both generation methods and detection strategies. It underscores why detection research must prioritise generalisation and resilience to evolving synthesis techniques, rather than relying on assumptions about specific models or controlled deployment scenarios (Cai et al., 2024; Dolhansky et al., 2020; Rossler et al., 2019). This perspective motivates the examination of deepfake generation techniques in the following

section, as well as the detection methodologies reviewed later in this chapter.

2.5 Deepfake Generation Techniques

Deepfake generation techniques can be broadly categorised according to the modality they manipulate: images, audio, or video. While early systems often focused on a single modality, recent advances increasingly integrate multiple streams, complicating forensic analysis. Understanding the mechanisms behind these generation techniques is essential for contextualising the strengths and limitations of current detection approaches (Cai et al., 2024; Dolhansky et al., 2020; Rossler et al., 2019; Yi et al., 2023).

2.5.1 Image-Based Manipulation

Image-based deepfake generation primarily targets facial appearance through identity replacement, expression transfer, or attribute manipulation (Rossler et al., 2019). Early approaches relied on autoencoder architectures trained to map facial representations between source and target identities, enabling face swapping with relatively limited training data (Dolhansky et al., 2020). These systems commonly employed a shared encoder with identity specific decoders, allowing facial expressions and pose information to be transferred while preserving identity related features (Güera and Delp, 2018).

Subsequent progress was driven by the adoption of Generative Adversarial Networks (GANs), which enabled higher resolution synthesis and more realistic texture generation (Rossler et al., 2019). Advances in training strategies and loss formulation, including stabilisation techniques such as the Two-Time Scale Update Rule, significantly reduced artefacts such as colour inconsistency and boundary distortion, making image-based deepfakes increasingly difficult to detect using handcrafted forensic cues (Heusel et al., 2017; Li et al., 2019). Recent studies have incorporated attention mechanisms and multi-scale discriminator architectures to improve performance under common post-processing operations such as compression and resizing, which can reduce visual differences between synthetic and authentic facial images (Cai et al., 2024).

2.5.2 Audio-Based Manipulation

Audio-based deepfake generation focuses on synthesising or converting speech to mimic the voice characteristics of a target speaker. Two dominant paradigms underpin these systems: text-to-speech (TTS), which generates speech directly from text, and voice

conversion (VC), which transforms the vocal attributes of a source speaker into those of a target speaker while preserving linguistic content (Yi et al., 2023). Early systems relied on statistical parametric models, but recent advances employ deep neural architectures capable of producing highly natural and expressive speech (Korshunov and Marcel, 2018). Modern audio deepfakes leverage end-to-end neural models, including encoder–decoder architectures and diffusion-based approaches, which allow high fidelity voice cloning from limited reference audio (Shen et al., 2023). These systems can accurately reproduce speaker identity, prosody, and emotional tone, making synthetic speech difficult to distinguish from genuine recordings by both humans and automated verification systems (Korshunov and Marcel, 2018; Yi et al., 2023). The increasing accessibility of such tools has lowered the barrier to misuse, particularly in impersonation and fraud scenarios, posing significant challenges for existing audio authentication mechanisms.

2.5.3 Video-Based Manipulation

Video-based deepfake generation extends beyond static image manipulation by modelling temporal consistency across frames, allowing for realistic facial motion, head pose variation, and synchronisation with speech (Rossler et al., 2019). Early video deepfakes often exhibited temporal artefacts such as flickering, inconsistent lighting, or unnatural motion, which could be exploited by detection systems (Li and Lyu, 2018). However, advances in spatio-temporal modelling have substantially mitigated these weaknesses.

Recent video-based techniques integrate facial reenactment, motion transfer, and lip synchronisation models to produce coherent and temporally stable outputs (Dolhansky et al., 2020). The emergence of multimodal generation frameworks further enables joint manipulation of video and audio streams, requiring generators to maintain consistency across visual appearance, speech content, and timing (Cai et al., 2024). Additionally, neural rendering approaches such as Neural Radiance Fields have been adopted to synthesise realistic talking heads with controllable viewpoints and lighting conditions, further increasing realism and generalisation across scenarios (Guo et al., 2021). These developments significantly complicate forensic analysis, as detectors must now take into account both spatial and temporal cues across multiple modalities.

2.6 Deepfake Detection Methodologies

Deepfake detection methodologies aim to distinguish synthetic media from authentic recordings by identifying artefacts, inconsistencies, or statistical patterns introduced during the generation process. As deepfake generation techniques have evolved, detection approaches have progressed similarly from handcrafted forensic features to data driven deep learning models. However, despite substantial advances, existing detection systems continue to struggle with generalisation, particularly under real world conditions (Dolhansky et al., 2020; Rossler et al., 2019).

2.6.1 Image-Based Detection Approaches

Image-based deepfake detection methods focus on identifying spatial inconsistencies within individual frames. Early approaches relied on handcrafted forensic cues, exploiting artefacts such as abnormal eye blinking, colour mismatches, and unnatural facial boundaries that were common in early face-swapping systems (Li and Lyu, 2018). While effective against early deepfakes, these methods lacked robustness as generative models improved. The adoption of convolutional neural networks (CNNs) marked a shift towards learning discriminative features directly from data. Large-scale benchmarks such as FaceForensics++ demonstrated that CNN-based detectors could achieve high accuracy within controlled datasets, often exceeding 0.95 AUC (Rossler et al., 2019). Subsequent research explored frequency-domain representations, showing that GAN-generated images exhibit characteristic spectral artefacts that can be exploited for detection (Li et al., 2019).

Despite strong performance in the dataset, image-based detectors exhibit significant performance degradation when evaluated in unseen datasets or under post processing operations such as compression and resizing (Dolhansky et al., 2020). This suggests that many models overfit to dataset specific artefacts rather than learning generator invariant features. Recent work incorporating attention mechanisms and multiscale feature extraction has shown incremental robustness improvements, but image only detection remains insufficient against modern, high quality deepfakes (Cai et al., 2024).

2.6.2 Audio-Based Detection Approaches

Audio-based detection methods aim to identify synthetic speech generated through text-to-speech or voice conversion systems. Early research adapted techniques from au-

tomatic speaker verification, using features such as Mel frequency cepstral coefficients and phase based representations to capture artefacts introduced by speech synthesis (Korshunov and Marcel, 2018). These studies demonstrated that synthetic speech often contains over smoothed spectral patterns and unnatural phase information.

More recent approaches employ deep neural architectures to learn discriminative representations directly from raw or minimally processed audio (Yi et al., 2023). While these models achieve strong benchmark performance, they remain vulnerable to unseen synthesis techniques and domain shifts, mirroring the limitations observed in image-based detection (Korshunov and Marcel, 2018). Furthermore, empirical studies show that human listeners struggle to reliably distinguish real and synthetic speech, reinforcing the need for automated audio deepfake detection systems (Yi et al., 2023).

2.6.3 Video Based Detection Approaches

Video-based detection approaches extend image level analysis by incorporating temporal information across frames. Early methods exploited temporal artefacts such as inconsistent head motion, unnatural blinking patterns, and frame-to-frame discontinuities (Li and Lyu, 2018). These cues enabled sequence level classification using temporal aggregation or recurrent architectures.

As generative models improved, many temporal artefacts became less pronounced. Contemporary video detectors therefore rely on spatio-temporal architectures, including 3D convolutional networks and attention-based temporal modelling, to capture subtle motion inconsistencies (Dolhansky et al., 2020; Rossler et al., 2019). Although these methods outperform frame-based detectors in controlled settings, they remain sensitive to compression, frame rate variation, and domain shifts common in real world video content (Dolhansky et al., 2020).

2.6.4 Multimodal Detection Approaches

The emergence of deepfakes that manipulate both audio and visual streams has driven increasing interest in multimodal detection approaches. These methods jointly analyse facial motion, lip synchronisation, and speech content to identify cross modal inconsistencies that are difficult for generative models to reproduce perfectly (Yi et al., 2023). Datasets such as FakeAVCeleb and AV-Deepfake1M have enabled systematic evaluation of multimodal detectors under coordinated manipulation scenarios (Cai et al.,

2024).

Multimodal detection systems generally demonstrate improved robustness compared to unimodal approaches, particularly in cross-dataset evaluations (Cai et al., 2024). However, their performance remains constrained by dataset bias, limited availability of synchronised training data, and increased computational complexity, which may hinder real-time deployment.

2.6.5 Generalisation and Cross Dataset Performance

A central challenge across all detection modalities is poor generalisation to unseen generators and real-world media conditions. Numerous studies report substantial performance drops when detectors trained on one dataset are evaluated on others generated using different synthesis methods (Dolhansky et al., 2020; Rossler et al., 2019). These findings indicate that many detectors rely on superficial cues rather than fundamental properties of synthetic media.

Cross-dataset evaluation has therefore become a critical benchmark for assessing practical effectiveness. Recent work emphasises that improving generalisation is more important than achieving marginal gains in closed set benchmarks (Cai et al., 2024).

2.7 Gaps in the Literature

Prior work shows that unimodal deepfake detectors often fail to generalise to unseen data, motivating increased interest in audio–visual detection methods (Cai et al., 2024; Dolhansky et al., 2020; Rossler et al., 2019). However, existing multimodal approaches exhibit several unresolved limitations.

Most audio–visual detectors rely on simple feature fusion, typically concatenating independently learned audio and visual representations (Yi et al., 2023). This design limits their ability to capture temporal relationships between speech and facial motion, reducing sensitivity to cross modal inconsistencies.

In addition, many systems remain vision centric, treating audio as a supplementary rather than an equal source of evidence (Cai et al., 2024; Rossler et al., 2019). As a result, the contribution of audio information is rarely quantified through systematic ablation or controlled evaluation.

Finally, current methods commonly assume that both modalities are reliable, although real world media often exhibit asymmetric degradation due to noise or compression (Dolhansky et al., 2020).

2.8 Conclusion

This chapter reviewed deepfake generation and detection techniques across image, audio, video, and multimodal domains. Advances in generative models have led to increasingly realistic synthetic media, reducing the effectiveness of early detection approaches based on visible artefacts (Dolhansky et al., 2020; Li et al., 2019; Rossler et al., 2019).

Across modalities, a consistent limitation is reduced performance outside controlled benchmarks, particularly when detectors encounter unseen generators or real world media conditions (Cai et al., 2024; Dolhansky et al., 2020; Rossler et al., 2019). Multimodal detection offers potential advantages, but remains constrained by design assumptions, limited analysis of modality contributions, and sensitivity to domain shift (Cai et al., 2024; Yi et al., 2023).

These findings highlight the need for detection approaches that prioritise generalisation across modalities and manipulation techniques, motivating the research direction adopted in this dissertation.

Chapter 3: Methodology

3.1 Introduction

This chapter outlines the approach taken to design, develop, and assess a deepfake detection system for audio-visual content utilizing the AV-Deepfake1M++ dataset (Cai et al., 2025). The aim is to meet research goals systematically while respecting hardware, storage, and computational limits.

The work targets content-driven deepfakes, where a real identity is kept but spoken content is altered via audio synthesis and lip synchronisation. Such changes are subtle and localised, making unimodal detection unreliable (Cai et al., 2024; Korshunov and Marcel, 2018). Therefore, a multimodal strategy is adopted that exploits the mismatches between speech and lip motion.

3.2 Research and Development Approach

This study uses a quantitative research approach based on secondary analysis of the AV-Deepfake1M++ dataset (Cai et al., 2024). Using an existing dataset avoids the cost and practical difficulties of collecting a large-scale audio-visual corpus, which would not be realistic within the scope of this project.

To manage development and experimentation, a simple Kanban-based workflow is used. Tasks such as data loading, preprocessing, feature extraction, model integration, and evaluation are treated as separate stages and worked on incrementally. This makes it easier to test individual parts of the system without having to run the full pipeline each time.

Each component is first tested using a small subset of video clips to confirm that it behaves as expected. Only after basic performance checks are satisfied is the component integrated into the full pipeline. This step-by-step approach helps reduce errors and unnecessary computation, which is particularly important when working with high-dimensional audio and video data.

3.3 System Architecture and implementation

3.3.1 Overview of the Proposed System

The main objective of this project is to maximise classification accuracy at the clip level. Although, the AV-Deepfake1M++ dataset provides temporal localisation annotations, these are not utilised in the current work due to the high computational cost of frame-level modelling. Given the available hardware and storage constraints, the study focuses on video-level classification.

3.3.2 Dataset Selection and Subset Strategy

The AV-Deepfake1M++ dataset is approximately 1.4 TB in size, which makes full local storage and processing infeasible within the available development environment. To manage this constraint, the study adopts a subset-based strategy that uses selected portions of the training and validation splits. The training subset is used to learn model parameters, while the validation subset supports hyper parameter tuning and early stopping. The TestA and TestB splits are not used during model development.

3.3.3 Audio Feature Extraction Module

For the audio modality, the system uses Wav2Vec 2.0, which is a pretrained self-supervised speech model that extracts feature embeddings directly from raw audio signals (Baevski et al., 2020). Unlike traditional handcrafted features such as MFCCs, Wav2Vec 2.0 learns contextual and prosodic information from speech, making it more suitable for detecting artefacts introduced by modern speech synthesis and voice cloning methods (Yi et al., 2023). Audio is extracted from each video clip and processed separately from the visual stream. The resulting embeddings capture temporal speech characteristics and are forwarded to the fusion module. No explicit temporal localisation is performed at this stage, as the focus of this study is on clip-level classification rather than frame-level analysis.

3.3.4 Visual Feature Extraction Module

The visual stream focuses on the mouth ROI to reduce computational overhead and remove irrelevant facial information. Facial landmark detection is used to locate and crop the lip region of each frame. Visual features are extracted using a lightweight con-

volutional neural network, such as MobileNetV3, which provides a balance between efficiency and representational capacity (Howard et al., 2019). The extracted features encode spatial characteristics related to lip shape and movement, which are particularly relevant for detecting lip-sync inconsistencies and mouth-related artefacts commonly observed in deepfake videos (Li and Lyu, 2018).

3.3.5 Cross-Modal Fusion and Classification

Audio and visual embeddings are combined using a cross-modal fusion module inspired by the DiMoDif architecture (Discourse Modality-information Differentiation) proposed for the AV-Deepfake1M++ dataset (Cai et al., 2025). DiMoDif is designed to model fine-grained dependencies between audio and visual streams by learning how speech content (phonemes) aligns with corresponding mouth movements (visemes) over time. This design directly targets content-driven deepfakes, allowing subtle inconsistencies introduced by manipulation to be detected.

In this study, the fusion module operates on temporally aligned audio and visual embeddings to learn a shared representation that captures cross-modal coherence. The resulting fused representation is passed to a binary classification head, which outputs the probability that a video being real or manipulated. The temporal boundary detection components of the original DiMoDif framework are not included, as the focus of this project is video-level classification.

3.4 Model Training and Evaluation Strategy

Training is performed on the selected training subset, with early stopping applied based on validation performance to reduce overfitting. The model is optimised using binary cross-entropy loss. Detection accuracy is used as the primary evaluation metric, with precision, recall, and F1-score reported to provide additional insight into classification performance. The TestA and TestB splits are not used for this project.

3.5 Challenges and Limitations

The primary limitation of this study arises from computational and storage constraints, which require the use of dataset subsets rather than full-scale training. Although this limits the model's exposure to the full diversity of manipulations, it allows for mean-

ingful experimentation within realistic hardware boundaries. Another limitation is the exclusion of temporal localisation. Although frame-level manipulation detection could provide deeper insight into model behaviour, such methods require significantly greater computational resources and are therefore deferred to future work. Finally, since the dataset is constructed by third parties, the study relies on the accuracy and completeness of existing annotations. Any bias present in the dataset can influence the performance of the model.

3.6 Ethical Considerations

The dataset used in this study is publicly available and anonymised, and contains no personally identifiable information. As such, no direct ethical concerns arise from data collection. However, deepfake detection research has broader ethical implications. Although the goal of this work is defensive, the same technologies can be misused to improve synthetic media generation (Chesney and Citron, 2019; Westerlund, 2019). Therefore, this study positions its contribution strictly within the context of detection and harm mitigation, without deployment or real-world enforcement claims.

3.7 Conclusion

A lightweight multimodal pipeline combining Wav2Vec 2.0 and MobileNetV3 features via a DiMoDif fusion module will be trained and evaluated on a representative subset under hardware constraints. Implementation details and results are presented in the following chapter.

Chapter 4: Results/Findings

Depending on your topic, you may have variety of things to report here. But these are objective reporting without stating your opinion on the results you are getting. This will happen in the next chapter.

4.1 Source Code

Use the Code style for presenting code snippets within this report. Please include snippets, rather than 1,000s of line of code which should be placed in an appendix or, if very large, in a separate document. If you have no code, then other implementation artifacts (e.g. screenshots) can be placed here if not evidenced elsewhere.

```
static public void main(String[] args) {  
    try {  
        UIManager.setLookAndFeel(UIManager.getSystemLookAndFeelClassName());  
    }  
    catch(Exception e) {  
        e.printStackTrace();  
    }  
    new WelcomeApp();  
}
```

Chapter 5: Evaluation

Include an evaluation of both product and process. You need to be objective in evaluating your work. Make sure you have a reflection here on what worked and what did not work.

Chapter 6: Conclusions

Write the conclusion of your project as well as any information related to future work. Remember to reflect on the key findings, limitations of these findings, and future opportunities to extend this work. Maybe a student next year can take your project and extend it!

Appendix A: Media Content

If the dissertation has a DVD or pendrive attached to it, you will need a section which explains what is on the media (structure, files, data, etc.). This could be a table with filename and description.

Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

Appendix B: Installation Instructions

Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

Appendix C: User Manual

Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

This is the second paragraph. Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

And after the second paragraph follows the third paragraph. Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

After this fourth paragraph, we start a new paragraph sequence. Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font,

how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like “Huardest gefburn”? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

References

- Allyn, B. (2022), ‘A deepfake video showing volodymyr zelenskyy surrendering worries experts : Npr’.
URL: <https://www.npr.org/2022/03/16/1087062648/deepfake-video-zelenskyy-experts-war-manipulation-ukraine-russia>
- Baevski, A., Zhou, H., Mohamed, A. and Auli, M. (2020), ‘wav2vec 2.0: A framework for self-supervised learning of speech representations’.
URL: <https://arxiv.org/abs/2006.11477>
- Bragg, J. (2025), ‘Ai-generated videos, photos used in new virtual kidnapping scams, fbi warns’.
URL: <https://www.axios.com/2025/12/05/ai-videos-photos-scam-virtual-fbi-victim-ransom-kidnapping>
- Cai, Z., Ghosh, S., Adatia, A. P., Hayat, M., Dhall, A., Gedeon, T. and Stefanov, K. (2024), ‘Av-deepfake1m: A large-scale llm-driven audio-visual deepfake dataset’, *Proceedings of the 32nd ACM International Conference on Multimedia (MM ’24)*, October 28-November 1, 2024, Melbourne, VIC, Australia **1**.
URL: <https://doi.org/10.1145/3664647.3680795>
- Cai, Z., Kuckreja, K., Chuchra, A., Khan, M. H., Gedeon, T. and Dhall, A. (2025), ‘Av-deepfake1m++: A large-scale audio-visual deepfake benchmark with real-world perturbations’, *Proceedings of 33rd ACM International Conference on Multimedia (ACMMM ’25)* **1**.
URL: <https://deepfakes1m.github.io/2025>.
- Chesney, B. and Citron, D. (2019), ‘Deep fakes: A looming challenge for privacy, democracy, and national security’, *California Law Review* **107**, 1753–1820.
URL: <https://www.californialawreview.org/print/deep-fakes-a-looming-challenge-for-privacy-democracy-and-national-security>
- Clark, A. (2025), ‘Deepfake videos impersonating real doctors push false medical advice and treatments - cbs news’.
URL: <https://www.cbsnews.com/news/deepfake-videos-impersonating-real-doctors-push-false-medical-advice-treatments/>

- Dolhansky, B., Bitton, J., Pflaum, B., Lu, J., Howes, R., Wang, M. and Ferrer, C. C. (2020), 'The deepfake detection challenge (dfdc) dataset'.
URL: <https://arxiv.org/pdf/2006.07397>
- Goodfellow, I., Bengio, Y. and Courville, A. (2016), *Deep Learning*, MIT Press. <http://www.deeplearningbook.org>.
- Guo, Y., Chen, K., Liang, S., Liu, Y. J., Bao, H. and Zhang, J. (2021), 'Ad-nerf: Audio driven neural radiance fields for talking head synthesis', *Proceedings of the IEEE International Conference on Computer Vision* pp. 5764–5774.
URL: <https://arxiv.org/pdf/2103.11078>
- Güera, D. and Delp, E. J. (2018), Deepfake video detection using recurrent neural networks, in '2018 15th IEEE International Conference on Advanced Video and Signal Based Surveillance (AVSS)', pp. 1–6.
- He, Y., Gan, B., Chen, S., Zhou, Y., Yin, G., Song, L., Sheng, L., Shao, J. and Liu, Z. (2021), 'ForgeryNet: A versatile benchmark for comprehensive forgery analysis', *Proceedings of the IEEE Computer Society Conference on Computer Vision and Pattern Recognition* pp. 4358–4367.
URL: <https://arxiv.org/pdf/2103.05630>
- Heusel, M., Ramsauer, H., Unterthiner, T., Nessler, B. and Hochreiter, S. (2017), 'Gans trained by a two time-scale update rule converge to a local nash equilibrium', *Advances in Neural Information Processing Systems* **2017-December**, 6627–6638.
URL: <https://arxiv.org/pdf/1706.08500>
- Hornik, K., Stinchcombe, M. and White, H. (1989), 'Multilayer feedforward networks are universal approximators', *Neural Networks* **2**.
- Howard, A., Sandler, M., Chu, G., Chen, L.-C., Chen, B., Tan, M., Wang, W., Zhu, Y., Pang, R., Vasudevan, V., Le, Q. V. and Adam, H. (2019), 'Searching for mobilenetv3'.
URL: <https://arxiv.org/abs/1905.02244>
- ITVNews (2024), '“deepfake doctors” are fuelling health scams on social media, bmj warns | itv news'.
URL: <https://www.itv.com/news/2024-07-18/deepfake-doctors-are-fuelling-health-scams-on-social-media-bmj-warns>
- Korshunov, P. and Marcel, S. (2018), 'Speaker inconsistency detection in tampered video', *European Signal Processing Conference* **2018-September**, 2375–2379.

- Lecun, Y., Bengio, Y. and Hinton, G. (2015), 'Deep learning', *Nature* **521**, 436–444.
- Li, Y. and Lyu, S. (2018), 'Exposing deepfake videos by detecting face warping artifacts'.
URL: <https://arxiv.org/pdf/1811.00656>
- Li, Y., Yang, X., Sun, P., Qi, H. and Lyu, S. (2019), 'Celeb-df: A large-scale challenging dataset for deepfake forensics', *Proceedings of the IEEE Computer Society Conference on Computer Vision and Pattern Recognition* pp. 3204–3213.
URL: <https://arxiv.org/pdf/1909.12962>
- Milmo, D. (2024), 'Uk engineering firm arup falls victim to £20m deepfake scam | engineering | the guardian'.
URL: <https://www.theguardian.com/technology/article/2024/may/17/uk-engineering-arup-deepfake-scam-hong-kong-ai-video>
- Moore, A. (2025), "it was about degrading someone completely": the story of mr deepfakes – the world's most notorious ai porn site | technology | the guardian'.
URL: <https://www.theguardian.com/society/ng-interactive/2025/dec/05/it-was-about-degrading-someone-completely-the-story-of-mr-deepfakes-the-worlds-most-notorious-ai-porn-site>
- Narayan, K., Agarwal, H., Thakral, K., Mittal, S., Vatsa, M. and Singh, R. (2023), Df-platter: Multi-face heterogeneous deepfake dataset, in '2023 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)', pp. 9739–9748.
- Rossler, A., Cozzolino, D., Verdoliva, L., Riess, C., Thies, J. and Niessner, M. (2019), 'Faceforensics++: Learning to detect manipulated facial images', *Proceedings of the IEEE International Conference on Computer Vision* pp. 1–11.
URL: <https://arxiv.org/pdf/1901.08971>
- Shen, K., Ju, Z., Tan, X., Liu, Y., Leng, Y., He, L., Qin, T., Zhao, S. and Bian, J. (2023), 'NaturalSpeech 2: Latent diffusion models are natural and zero-shot speech and singing synthesizers'.
URL: <https://arxiv.org/abs/2304.09116>
- Stupp, C. (2019), 'Fraudsters used ai to mimic ceo's voice in unusual cybercrime case - wsj'.
URL: <https://www.wsj.com/articles/fraudsters-use-ai-to-mimic-ceos-voice-in-unusual-cybercrime-case-11567157402>

- Vaccari, C. and Chadwick, A. (2020), 'Deepfakes and disinformation: Exploring the impact of synthetic political video on deception, uncertainty, and trust in news', *Social Media and Society* **6**.
- Westerlund, M. (2019), 'The emergence of deepfake technology: A review'.
- Yi, J., Wang, C., Tao, J., Zhang, X., Zhang, C. Y. and Zhao, Y. (2023), 'Audio deepfake detection: A survey'.
- URL:** <https://arxiv.org/abs/2308.14970>
- Zhang, B., Cui, H., Nguyen, V. and Whitty, M. (2025), 'Audio deepfake detection: What has been achieved and what lies ahead', *Sensors (Basel, Switzerland)* **25**, 1989.
- URL:** <https://pmc.ncbi.nlm.nih.gov/articles/PMC11991371/>
- Zhou, T., Wang, W., Liang, Z. and Shen, J. (2021), 'Face forensics in the wild', *Proceedings of the IEEE Computer Society Conference on Computer Vision and Pattern Recognition* pp. 5774–5784.
- URL:** <https://arxiv.org/pdf/2103.16076>
- Zi, B., Chang, M., Chen, J., Ma, X. and Jiang, Y. G. (2021), 'Wilddeepfake: A challenging real-world dataset for deepfake detection', *MM 2020 - Proceedings of the 28th ACM International Conference on Multimedia* pp. 2382–2390.
- URL:** <https://arxiv.org/pdf/2101.01456>